




Recovery and Buffer Management (Log based recovery)

Quick Review

- Last time before a midterm, we have **log based recovery topic.**



- idea is quite simple, we keep log record.
 - When we start transaction, we have **<Tx, Start>**
 - When the transaction is active during processing of transaction, we have **<Tx, Q, ov, nv>** (old value, new value)
 - When the transaction is committed, we have **<Tx, Commit>** output to the log file.
 - So the physical commit (when we have log-based recovery) → when commit record output to the log file successfully.
- Based on this principle, During recovery time after failure, the system simply search commit record from the back
 - if the commit record is found → that transaction commit
 - otherwise → the transaction failed
- Even though, you have commit record is log buffer → it not yet commit (At failure time, the log record in log buffer maybe lose) → we count only commit in log file or commit in AIJ

👉 Old value/ new value

- we use new value to redo → solve first problem (commit before output)
- we used old value to undo DB space → solve second problem (output before commit)

Old value is store in BIJ, new value stored in AIJ → can choose to install one of them or both of them

Content

🌱 2 Approaches of log-based recovery

(Book has slight different terminology, same idea but different terminology)



Log-Based Recovery

- A **log** is a sequence of **log records**. The records keep information about update activities on the database.
 - The **log** is kept on stable storage
- When transaction T_i starts, it registers itself by writing a $\langle T_i \text{ start} \rangle$ log record
- Before T_i executes **write**(X), a log record $\langle T_i, X, V_1, V_2 \rangle$ is written, where V_1 is the value of X before the write (the **old value**), and V_2 is the value to be written to X (the **new value**).
- When T_i finishes its last statement, the log record $\langle T_i \text{ commit} \rangle$ is written.
- Two approaches using logs
 - Immediate database modification
 - Deferred database modification.



From book

- There are 2 approaches using log (log based recovery)

1. Immediate database modification

- The immediate-modification scheme allows update of an uncommitted transaction to be made to the buffer, or the disk itself, before the transaction commits
- Output may take place before or after commit



- in this categories

- BIJ only is an immediate DB modification (take place before commit)
 - no new value → must output before commit
- Both AIJ & BIJ is an immediate DB modification (output anytime)



Immediate Database Modification

- The **immediate-modification** scheme allows updates of an uncommitted transaction to be made to the buffer, or the disk itself, before the transaction commits
- Update log record must be written *before* database item is written
 - We assume that the log record is output directly to stable storage
 - (Will see later that how to postpone log record output to some extent)
- Output of updated blocks to disk can take place at any time before or after transaction commit
- Order in which blocks are output can be different from the order in which they are written.
- The **deferred-modification** scheme performs updates to buffer/disk only at the time of transaction commit
 - Simplifies some aspects of recovery
 - But has overhead of storing local copy



Immediate Database Modification Example

Log	Write	Output
$\langle T_0 \text{ start} \rangle$		
$\langle T_0, A, 1000, 950 \rangle$		
$\langle T_0, B, 2000, 2050 \rangle$		
	$A = 950$ $B = 2050$	
$\langle T_0 \text{ commit} \rangle$		
$\langle T_1 \text{ start} \rangle$		
$\langle T_1, C, 700, 600 \rangle$		
	$C = 600$	
$\langle T_1 \text{ commit} \rangle$		
		B_B, B_C
		B_A

B_C output before T_1 commits

B_A output after T_0 commits

- Note: B_X denotes block containing X.



Immediate DB Modification Recovery Example

Below we show the log as it appears at three instances of time.

<T ₀ start>	<T ₀ start>	<T ₀ start>
<T ₀ , A, 1000, 950>	<T ₀ , A, 1000, 950>	<T ₀ , A, 1000, 950>
<T ₀ , B, 2000, 2050>	<T ₀ , B, 2000, 2050>	<T ₀ , B, 2000, 2050>
	<T ₀ commit>	<T ₀ commit>
	<T ₁ start>	<T ₁ start>
	<T ₁ , C, 700, 600>	<T ₁ , C, 700, 600>
		<T ₁ commit>
(a)	(b)	(c)

Recovery actions in each case above are:

- (a) undo (T₀): B is restored to 2000 and A to 1000.
- (b) undo (T₁) and redo (T₀): C is restored to 700, and then A and B are set to 950 and 2050 respectively.
- (c) redo (T₀) and redo (T₁): A and B are set to 950 and 2050 respectively. Then C is set to 600

- in immediate database modification, we need an old value, because immediate imply that output can be done before commit, so we need old value for recovery process (undo)
- this diagram show both old and new value

2. Deferred database modification (can be postpone)

- The deferred-modification scheme performs updates to buffer/disk only at the time of transaction commit
- It must commit first - commit then later output
- Output may take place after commit or at the commit time



AIJ only is deferred DB modification (we dont have old value, not allow db buffer to output to db space before commit)

Deferred Database Modification (Cont.)

- Below we show the log as it appears at three instances of time.

<T ₀ start>	<T ₀ start>	<T ₀ start>
<T ₀ , A, 950>	<T ₀ , A, 950>	<T ₀ , A, 950>
<T ₀ , B, 2050>	<T ₀ , B, 2050>	<T ₀ , B, 2050>
	<T ₀ commit>	<T ₀ commit>
	<T ₁ start>	<T ₁ start>
	<T ₁ , C, 600>	<T ₁ , C, 600>
		<T ₁ commit>
(a)	(b)	(c)


- If log on stable storage at time of crash is as in case:
 - No redo actions need to be taken
 - redo(T₀) must be performed since <T₀ commit> is present
 - redo(T₀) must be performed followed by redo(T₁) since <T₀ commit> and <T₁ commit> are present

- in this, it has only AIJ (new value)
- Transaction T0 transfer 50 dollars from Transaction A to B
 - Before the transaction A was 1000, B was 2000
 - After transaction A become 950, and B become 2050
 - The log only show new value (AIJ)
- For transaction T1, we withdraw 100 dollars from C (C was 700)
 - After withdrawal → become 600

Show only new value



Buffer Management (19.5 - book)



from, this we see log buffer, db buffer and know that they are in main memory

- Log buffer and DB buffer are in main memory
- Who manage the buffer (DB buffer and log buffer)? Should it be manage by the DBMS or the OS?
 - when we study OS → we know that OS controls the main memory. It can paging, segmentation,, It also has page replacement strategy like FIFO, read recently used, read frequently used, that adopt by the OS
 - OS actually control the main memory, but should we let OS handle buffer management of this log buffer and database buffer?
 - (There were discussion and also implementation.) How much involvement should DBMS be in this buffer management activity?
 - after long discussion and several implementation, it is very clear that it is the DBMS that should be in charge of this buffer management.

👉 There are protocol that the DBMS must execute and must guide, force DBMS to follow such protocol.

Those Protocol?

(Some topic from book)

1. **Log-Record Buffering** - Log record are buffered in main memory, instead of being output directly to stable storage and it need to be managed.



Log Record Buffering

- **Log record buffering:** log records are buffered in main memory, instead of being output directly to stable storage.
 - Log records are output to stable storage when a block of log records in the buffer is full, or a **log force** operation is executed.
 - Log force is performed to commit a transaction by forcing all its log records (including the commit record) to stable storage.
 - Several log records can thus be output using a single output operation, reducing the I/O cost.
- Log force is performed to commit a transaction by forcing all its log records (including the commit record) to stable storage.
 - This book always assume that log file is in stable storage → so it mean force to log file. (when we say force to stable storage)
 - Several log records can thus be output using a single output operation, reducing the I/O cost.



There an important protocol on **Log Record Buffering**

Log Record Buffering (Cont.)

- The rules below must be followed if log records are buffered:
 - Log records are output to stable storage in the order in which they are created.
 - Transaction T_i enters the commit state only when the log record $\langle T_i \text{ commit} \rangle$ has been output to stable storage.
 - Before a block of data in main memory is output to the database, all log records pertaining to data in that block must have been output to stable storage.
 - This rule is called the **write-ahead logging** or **WAL** rule
 - Strictly speaking, WAL only requires undo information to be output

Rule of Log-record buffering

- Log records are output to stable storage in the order in which they are created
 - this imply that commit record must be the last one of that transaction.
- Transaction T_i enters the commit state only when the log record $\langle T_i \text{ commit} \rangle$ as been output to stable storage. (to ensure physical commit)

- WAL → before block of data output to the database, log record output to stable storage first.



The DBMS must follow this rule


- For OS point of view, DBMS is just an application program for OS
 - So, this is application requirement of DBMS for log based recovery.
- It must be DBMS that instruct the OS that ok, you must go now or wait → they must be test together.

2. Database Buffering



Database Buffering


- Database maintains an in-memory buffer of data blocks
 - When a new block is needed, if buffer is full an existing block needs to be removed from buffer
 - If the block chosen for removal has been updated, it must be output to disk
- The recovery algorithm supports the **no-force policy**: i.e., updated blocks need not be written to disk when transaction commits
 - force policy**: requires updated blocks to be written at commit
 - More expensive commit
- The recovery algorithm supports the **steal policy**: i.e., blocks containing updates of uncommitted transactions can be written to disk, even before the transaction commits



- Suppose you have block b1 in DB buffer and a block b2 in DB space → in practice, after a while your database buffer gonna be full because you use data (input data to buffer), and if you want b2 to be input in DB buffer, you can't just input b2 directly to buffer → b2 has to replace a block here (DB space is full). → bring b1 out before you move b2 in
- Step (database buffer simple protocol)
 - When you bring b1 out, you need to output the log of b1 first to the log file
 - b1 back to DB space
 - b2 in to db buffer
- DB buffer also requires the WAL protocol

3. Operating System role in buffer management

- We can manage the database buffer by using one of two approaches:
1. The database system reserves part of main memory to serve as a buffer that it (DBMS), rather than the operating system, manages.
 - DBMS reserves part of main memory to serve as a buffer that the DBMS itself manages



- From main memory pov, the os allow the DBMS to reserves part of the main memory as database and log buffer
 - size of database and log buffer is fixed size → The DBMS occupied a fixed size of main memory and it's managed by the DBMS.
 - It's the DBMS that manage the protocol without passing through the OS.
 - **blue part → virtual memory**
2. The database system implements its buffer within the virtual memory provided by the operating system.
 - it is the DBMS that implement it buffer within the virtual memory provided by the OS.

Diagram show interaction between DBMS and OS.

